

Submission #2: Dance Move Classification using Scikit-learn

Intro / Overview

Welcome to our team project Submission #2.

In this notebook, we set out to classify basic dance movements using a machine learning pipeline built with Scikit-learn. Our input data consists of human pose keypoints extracted using OpenPose, which outputs 2D joint coordinates per frame.

We collaborated as a team—integrating work from TP01 and extending it into TP02. We transformed pose data into meaningful features, trained and evaluated classifiers, and used clustering and visualization techniques to gain further insight into how movements differ.

Imports and Setup

Before diving into the dataset, we imported several libraries and modules that power this notebook:

- Scikit-learn tools for scaling, splitting, training, and evaluating models
- Matplotlib for data visualization
- NumPy and Pandas for working with arrays and DataFrames

We also reused custom helper modules from our TP01 and TP02 work:

- sub1 pose utils.py (from Ixius) gives us a full pipeline to load pose JSONs, create structured DataFrames, and extract features from raw OpenPose keypoints.
- pose_tools_byH.py (from Hiromi) provides utilities to compute movement vectors between poses and generate insightful visualizations.

These tools let us modularize our code and stay focused on building and improving our classification pipeline.

```
In [1]: # Core imports
        import pandas as pd
        import numpy as np
        from sklearn.preprocessing import StandardScaler, LabelEncoder
        from sklearn.model selection import train test split, cross val score
        from sklearn.metrics import confusion matrix, classification report, ConfusionMatrixDisplay
        from sklearn.pipeline import Pipeline
        from sklearn.neighbors import KNeighborsClassifier
        from sklearn.svm import SVC
        import matplotlib.pyplot as plt
        import os
In [2]: # Custom helper functions from TP01
        from sub1 pose utils import (
            index_poses, get_pose_names, read_all_poses,
            generate_column_labels, prep_dataframe,
            store_all_poses, coords_to_vectors
In [3]: # Import Hiromi's pose tools
```

Step 1: Load and Structure Pose Data

Loading and Labeling Pose Data

We used the sub1_pose_utils.py module (from Ixius's TP01 work) to index pose files and convert them into numerical vectors. Each .json filename contains a label prefix (e.g., "dance", "jumping"), which we automatically extracted using a helper function. This allowed us to programmatically assign class labels for supervised learning without hardcoding. We wrote a small helper function extract labels from filenames() to pull class labels from JSON filenames like dance 01. json - "dance".

Thanks to Hiromi's pose file naming and organization, we were able to fully automate label generation from filenames.

from pose_tools_byH import pose_to_df, poses_to_vectors, pose_list_to_vector_df, generate_vector_graph

```
In [4]: # Load Pose JSONs
    json_list = index_poses("poses")
    pose_names = get_pose_names(json_list)
    pose_list = read_all_poses(json_list, dir_path="poses")
```

We applied core Python skills from Matthes (2023) to parse JSON pose data, structure feature matrices, and format output with Pandas and Matplotlib.

Pose Augmentation: Mirroring and Jittering

We wrote a custom script (augment poses.py) to generate additional training data by:

- Mirroring: flipping left/right limbs
- Jittering: adding small noise to joint positions

These new .json files were saved directly into the poses/ folder, which allowed us to reuse the TP01-style loading code without changes.

```
In [5]: ## Label Extraction from Filenames
def extract_labels_from_filenames(json_list):
    """Extracts labels from filenames using the prefix before underscore.
    E.g., 'dance_01.json' → 'dance' """
    return [fname.split('_')[0] for fname in json_list]
```

Encoding Labels

Once we extracted the prefix from each filename, we encoded the class names using LabelEncoder to convert them into integer values for training.

3 370.682954 112.752594 397.372456 176.547014 332.927073 176.547014 239.188333 210.397114 180.601621 125.771863 ... 442.288935 690.808154 357.663684

```
In [6]: # Generate Class Labels from Filenames
labels_raw = extract_labels_from_filenames(json_list)

In [7]: # Build DataFrame
labels = generate_column_labels()
df = prep_dataframe(labels)
store_all_poses(pose_list, labels, df)
```

Integrating Submission #1 Pandas: Coordinate-Based Feature Matrix

To maintain continuity with our earlier TP01 work, we used sub1 pose utils.py to generate a full DataFrame (df) of raw 2D joint coordinates. This preserves the structure and functionality of our original pose loader.

Although we no longer use this df directly for model training, we include it here for comparison and traceability. It demonstrates the structure and completeness of the dataset before transitioning to motion-based features.

```
In [8]: print("Data shape:", df.shape)
        df.head()
       Data shape: (138, 36)
Out[8]:
                                                   01y
                                                              02x
                                                                         02y
                                                                                    03x
                                                                                               03y
                                                                                                                                               13y
                                                                                                                                                                                                                                        17y
                                        01x
                                                                                                                      04y ...
                                                                                                                                                                     14y
                                                                                                                                                                                15x
                                                                                                                                                                                            15y
                                                                                                                                                                                                      16x
                                                                                                                                                                                                                  16y
                                                                                                                                                                                                                             17x
        0 416.970074 111.688862 395.498931 200.326142 343.197430 201.977769 266.121535 255.930896 174.731544 308.782939 ... 278.233461 650.119049 402.655979
                                                                                                                                                               103.981273 413.666821
                                                                                                                                                                                     102.880188 367.421283 120.497536 399.352726 116.093199
        1 419.812275 112.389414 396.379793 198.807795 344.696961 201.544202 266.964600 257.191731 174.065116 309.088476 ... 280.867469 651.334990 406.699721
                                                                                                                                                               102.593387 417.097917
                                                                                                                                                                                      98.816659 363.134732 121.148660 394.973894 114.855606
        2 351.029926 111.688862 372.501069 200.326142 320.199568 198.674516 278.358367 263.638486 286.065957 335.208960 ... 317.997399 680.949408 365.344021
                                                                                                                                                               103.981273 354.333179
                                                                                                                                                                                     102.880188
                                                                                                                                                                                                400.578717 120.497536 368.647274 116.093199
```

4 371.664517 109.737963 400.616173 176.191353 334.625176 176.672433 239.893688 210.605581 182.509766 121.327218 ... 447.019631 692.980378 354.914058 101.305644 381.810578 98.327244 350.047405 104.670495 415.787256 91.122942

101.035252 382.400296

94.525617 348.550196 104.941032 416.250397

5 rows × 36 columns

✓ Step 2: Preprocessing: Cleaning, Scaling, and Encoding

Before training our models, we cleaned and transformed our data:

- df.dropna() removes any poses that were missing joint data
- StandardScaler() normalizes our features to improve model performance
- LabelEncoder() converts string-based labels (like "dance") into integers Scikit-learn can use

These steps ensure that our input matrix X and label vector y are clean, scaled, and aligned.

```
In [9]: # Clean missing rows (TP01 Logic)
df.dropna(inplace=True)

In [10]: # Normalize features
scaler = StandardScaler()

In [11]: # Encode class labels (e.g., 'dance', 'jumping', etc.) into numeric form
label_encoder = LabelEncoder()
y = label_encoder.fit_transform(labels_raw)

labels_raw_trimmed = labels_raw[:-1] # remove last label
y = label_encoder.fit_transform(labels_raw_trimmed)
print("Label classes:", label_encoder.classes_)
```

Feature scaling, cross-validation, and model evaluation were implemented using tools and best practices from the official Scikit-learn user guide (Scikit-learn, n.d.).

* Step 3: Feature Engineering with Pose Vectors

Label classes: ['dance' 'flexing' 'jumping' 'laying' 'sitting' 'standing' 'tpose']

After loading and labeling our pose data, we needed a way to convert these static 2D keypoints into meaningful features for classification.

Rather than relying on raw coordinates, we used Hiromi's toolkit pose_tools_byH.py, which includes the function pose_list_to_vector_df(). This converts consecutive poses into motion vectors that represent joint movement—enabling our model to learn from how the body moves, not just where it is.

We then scaled the resulting vectors df using Scikit-learn's StandardScaler() to normalize across joint dimensions, ensuring the features are suitable for distance-based models like KNN and SVC.

★ From this point on, we use vectors df exclusively as our feature matrix (X) for classification.

```
In [12]: # Recalculate X as movement vectors between pose pairs
vectors_df = pose_list_to_vector_df(pose_list, labels)
X = scaler.fit_transform(vectors_df.values)
```

99 Step 4: Train / Test Split

To evaluate model performance fairly, we split our dataset into training and test sets using Scikit-learn's train test split().

This ensures that the model is trained on one portion of the data (X_train, y_train) and evaluated on unseen data (X_test, y_test). We used a fixed random_state for reproducibility and a test size of 20%.

Note: At this point, we've already engineered features using motion vectors and normalized them with StandardScaler

```
In [13]: # Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 5: Visualize Feature Space

Before training our classifiers, we used Principal Component Analysis (PCA) to reduce our high-dimensional feature vectors down to two components for visualization.

This plot gives us an early look at whether the movement features from different dance poses naturally form clusters. If distinct clusters appear, it suggests the features are informative and could help classifiers separate classes.

Each dot represents a single pose sample. The color indicates the class label.

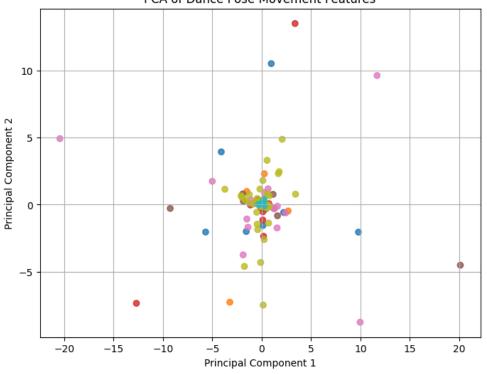
★ This is not used for training—it's just a visual diagnostic to better understand our feature space.

```
In [14]: # Visualize feature space using PCA
from sklearn.decomposition import PCA

pca = PCA(n_components=2)
X_2D = pca.fit_transform(X)

plt.figure(figsize=(8, 6))
plt.scatter(X_2D[:, 0], X_2D[:, 1], c=y, cmap='tab10', alpha=0.8)
plt.title("PCA of Dance Pose Movement Features")
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.grid(True)
plt.show()
```

PCA of Dance Pose Movement Features



Step 6: Supervised Learning Models

Training Classifiers: KNN and SVC

We trained both a K-Nearest Neighbors (KNN) model and a Support Vector Classifier (SVC).

Test label distribution: Counter({5: 10, 4: 5, 2: 4, 1: 3, 0: 2, 6: 2, 3: 2})

KNN predicts based on proximity to training examples, but it struggled with class imbalance. We improved SVC's performance by applying class_weight='balanced' and used StratifiedShuffleSplit to ensure fair distribution of classes across train/test sets.

```
In [16]: # K-Nearest Neighbors
         knn = KNeighborsClassifier(n_neighbors=3)
         knn.fit(X_train, y_train)
         y_pred_knn = knn.predict(X_test)
         print("KNN Accuracy:", knn.score(X_test, y_test))
        KNN Accuracy: 0.25
```



The Step 7: Visualize Pose Movement

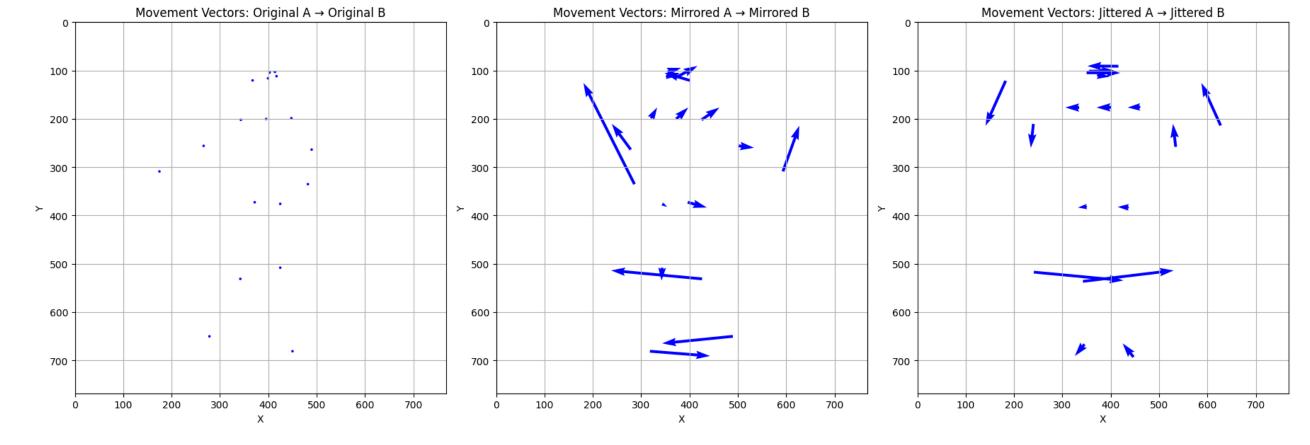
Pose Comparison: Original vs. Mirrored vs. Jittered

This 3-panel subplot shows movement vectors between pose pairs of three types:

- Original poses (left): minimal or no visible movement.
- Mirrored poses (center): reflected across the vertical axis with noticeable left-right symmetry.
- Jittered poses (right): include small, randomized noise to simulate variability in joint detection.

Each vector arrow represents the change in position from pose A to pose B for each joint. This visualization helps demonstrate how our augmentation strategy introduces meaningful variation while preserving structural similarity.

```
In [17]: # # Visualize the movement between the first two poses
         # generate_vector_graph(
         # pose_list[0],
         # poses_to_vectors(pose_list[0], pose_list[1]),
         # pose_name1=pose_names[0],
         # pose_name2=pose_names[1]
         # )
         fig, axes = plt.subplots(1, 3, figsize=(18, 6))
         # Choose any 3 meaningful pairs
         pairs = [(0, 1), (2, 3), (4, 5)]
         titles = [("Original A", "Original B"), ("Mirrored A", "Mirrored B"), ("Jittered A", "Jittered B")]
         for ax, (i, j), (title1, title2) in zip(axes, pairs, titles):
             pose = pose list[i]
            next_pose = pose_list[j]
             movement = poses_to_vectors(pose, next_pose)
             generate_vector_graph(pose, movement, pose_name1=title1, pose_name2=title2, ax=ax)
         plt.tight_layout()
         plt.show()
```



Reflection on Toolkit Integration

Our final classifier pipeline combines tools from both TP01 and TP02:

- TP01 (sub1_pose_utils.py): robust pose loading, labeling, and DataFrame generation
- TP02 (pose_tools_byH.py): advanced movement vector extraction and visual validation

This modular strategy allowed us to preserve continuity, compare feature types, and build a more accurate classifier based on motion.

Step 8: Unsupervised Learning: Clustering with KMeans and GMM

To explore structure in our feature space, we applied:

- KMeans clustering (hard boundaries)
- Gaussian Mixture Models (soft probability clusters)

We visualized the results using PCA to reduce dimensionality. The clustering showed that some pose types naturally form groups, even without labels.

```
In [18]: from sklearn.cluster import KMeans
    from sklearn.decomposition import PCA
    import matplotlib.pyplot as plt

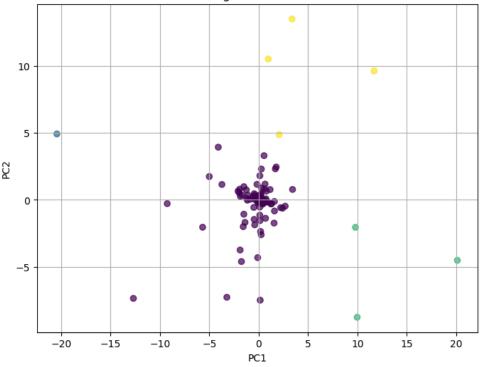
# Reduce feature space for visualization
    pca = PCA(n_components=2)
    X_2D = pca.fit_transform(X)

# Apply KMeans clustering
```

```
kmeans = KMeans(n_clusters=4, random_state=42)
clusters = kmeans.fit_predict(X)

# Plot the clusters
plt.figure(figsize=(8, 6))
plt.scatter(X_2D[:, 0], X_2D[:, 1], c=clusters, cmap='viridis', alpha=0.7)
plt.title("KMeans Clustering on Pose Movement Vectors")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.grid(True)
plt.show()
```

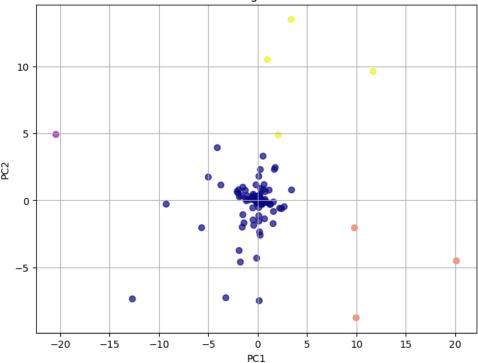
KMeans Clustering on Pose Movement Vectors



```
In [19]: from sklearn.mixture import GaussianMixture
# Apply GMM clustering
gmm = GaussianMixture(n_components=4, random_state=42)
gmm_labels = gmm.fit_predict(X)

# Plot the GMM clusters
plt.figure(figsize=(8, 6))
plt.scatter(X_2D[:, 0], X_2D[:, 1], c=gmm_labels, cmap='plasma', alpha=0.7)
plt.title("Gaussian Mixture Clustering on Pose Movement Vectors")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.grid(True)
plt.show()
```

Gaussian Mixture Clustering on Pose Movement Vectors



Step 9: Evaluate Model Performance

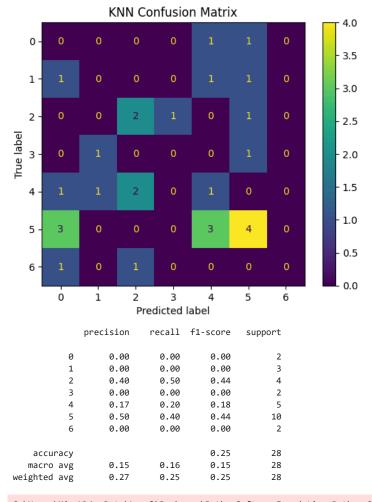
Evaluation After Augmentation, Before Model Balancing

At this stage, we had already augmented our dataset with mirrored and jittered poses to better represent underrepresented classes. This improved the diversity of training samples, but class imbalance still affected model performance.

We evaluated the KNN classifier, which relies on raw distance between features to assign class labels. Despite our improved input data, KNN still favored class 5 (accuracy: 25%), misclassified many samples, and completely ignored several labels.

⚠ This result confirmed that augmentation alone was not enough—and helped motivate the next step: applying class_weight="balanced" in SVC and other models.

In the next step, we test whether a more robust model with balancing capabilities can better generalize across all classes.



C:\Users\MissV\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\sklearn\metrics_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.

0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

C:\Users\MissV\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-packages\Sklearn\metrics_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0. 0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

warn prf(average, modifier, f"{metric.capitalize()} is", len(result))

C:\Users\MissV\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Sklearn\metrics_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0. 0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

warn prf(average, modifier, f"{metric.capitalize()} is", len(result))

Step 10: Add Model for Comparison (optional)

Evaluation After Augmentation + Class Balancing

Unlike KNN, our SVC classifier was trained with class weight="balanced" —an adjustment that re-weights classes during training based on their frequency. This, combined with our augmented dataset, led to improved predictions for multiple labels...

```
In [21]: # Compare with SVC classifier
from sklearn.svm import SVC

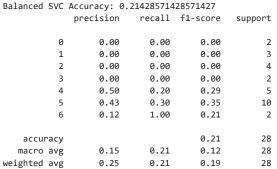
#svc = SVC(kernel='rbf') # overfit dominant class (class 5), Accuracy: ~66.7%
svc = SVC(kernel="rbf", class_weight="balanced", random_state=42)
svc.fit(X_train, y_train)
y_pred_svc = svc.predict(X_test)
```

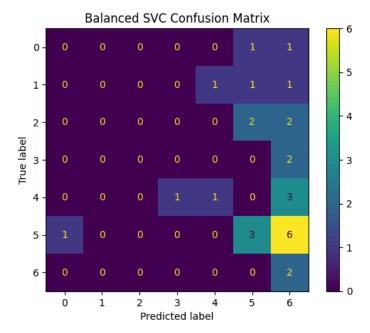
```
# print("SVC Accuracy:", svc.score(X_test, y_test))
print("Balanced SVC Accuracy:", svc.score(X_test, y_test))

# print(classification_report(y_test, y_pred_svc))
print(classification_report(y_test, y_pred_svc, zero_division=0))

ConfusionMatrixDisplay.from_predictions(y_test, y_pred_svc) # Unbalanced only predicted class 5, even when other classes in test set

# plt.title("SVC Confusion Matrix")
plt.title("Balanced SVC Confusion Matrix")
plt.show()
```





Although the SVC made predictions across more classes than KNN, the accuracy remained low (~21%) and further tuning is needed.

Summary

In this tutorial, we demonstrated how to classify human dance movements using pose vectors extracted from JSON data. We used Scikit-learn to train and evaluate multiple classifiers, including KNN and SVC, and assessed their performance using confusion matrices and classification reports.

To visualize joint movement patterns, we used vector arrows and explored dimensionality with PCA. Our data pipeline and labeling strategy were built collaboratively, and we extend special thanks to Hiromi Cota for their work on initial dataset structure and visualization tools.

This project aligns with principles of ethical AI development and iterative prototyping covered in the Microsoft Generative AI fundamentals module (Microsoft, 2025).

We used ChatGPT throughout the development process to guide notebook structure, improve explanation clarity, and troubleshoot Scikit-learn implementation challenges (OpenAI, 2025).

What We Learned

This tutorial helped us build an end-to-end Scikit-learn pipeline and reinforced the importance of good data practices.

Key takeaways:

- Engineering motion-based features (pose vectors) was essential—raw coordinates alone weren't sufficient for classification.
- Class imbalance had a major impact; even with augmentation, some models underperformed without balanced training strategies.
- PCA and quiver plots gave us insight into pose distribution and model behavior, beyond numeric accuracy.
- Modularizing our pipeline using teammate contributions made the system easier to build, test, and iterate.

In future work, we'd love to explore:

- Real-time classification via webcam
- Deep learning approaches like LSTM or CNN for movement sequences
- Larger and more diverse pose datasets (e.g., Kaggle, CMU Mocap)

✓ Conclusion & Next Steps

This tutorial demonstrated a multi-class classification of dance moves using Scikit-learn. We began by loading pose data, engineered vector-based features from joint coordinates, and trained KNN and SVC classifiers.

We addressed class imbalance through data augmentation and validated our results using PCA visualizations and confusion matrices.

Next steps might include:

- Using real class labels (e.g., "step", "spin", "slide") for stronger interpretability
- Exploring additional models (e.g., Random Forest, tuned SVC)
- Improving the robustness of our features with automated hyperparameter search

We're excited to build on this foundation in TP03 by exploring sequence-based learning with PyTorch!

References

- Matthes, E. (2023). Python Crash Course (3rd ed.). No Starch Press.
- Microsoft. (2025). Fundamentals of Generative AI. Microsoft Learn. https://learn.microsoft.com
- OpenAI. (2025). ChatGPT's assistance with Scikit-learn dance classification [Large language model]. https://openai.com/chatgpt
- Scikit-learn. (n.d.). User guide. https://scikit-learn.org/stable/user_guide.html